

```
import pandas as pd
import numpy as np

# Load the CSV files into DataFrames
df1 = pd.read_csv('/content/indian_cities - kaggle dataset (2).csv')
df2 = pd.read_csv('/content/total-final-2 (1).csv')
unique_values=df2['city'].value_counts()

# Print or use the unique values as needed
print("Unique values in 'column_name':", unique_values)
# Merge the DataFrames based on a common column
merged_df = pd.merge(df1, df2, on='city', how='inner')

# # Save the merged DataFrame to a new CSV file
# merged_df.to_csv('merged_file.csv', index=False)
```

```
-----
FileNotFoundError                                Traceback (most recent call last)
```

```
<ipython-input-2-dafb614e6e2c> in <cell line: 5>()
```

```
3
```

```
4 # Load the CSV files into DataFrames
```

```
----> 5 df1 = pd.read_csv('/content/indian_cities - kaggle dataset (2).csv')
```

```
6 df2 = pd.read_csv('/content/total-final-2 (1).csv')
```

```
7 unique_values=df2['city'].value_counts()
```

```
⬆ 6 frames
```

```
/usr/local/lib/python3.10/dist-packages/pandas/io/common.py in get_handle(path_or_buf, mode, encoding,
compression, memory_map, is_text, errors, storage_options)
```

```
854     if ioargs.encoding and "b" not in ioargs.mode:
```

```
855         # Encoding
```

```
--> 856         handle = open(
```

```
857             handle,
```

```
858             ioargs.mode,
```

```
FileNotFoundError: [Errno 2] No such file or directory: '/content/indian_cities - kaggle dataset (2).csv'
```

```
import numpy as np
import pandas as pd
```

```
df= pd.read_csv("/content/Sheet 1-1-merged_file.xlsx - Sheet 1-1-merged_file (2).csv")
```

> 1. Basic Exploratory Data Analysis and inferences.

```
[ ] ↳ 7 cells hidden
```


✓ 2. Visualization using various plots and their inferences.

```
import matplotlib.pyplot as plt
import seaborn as sns
```

✓ a) Histogram


```
plt.hist(df['Rank'])
```

```
(array([3458., 455., 273., 0., 0., 0., 0., 0., 0.,
       91.]),
array([ 1., 49.5, 98., 146.5, 195., 243.5, 292., 340.5, 389.,
       437.5, 486. ]),
<BarContainer object of 10 artists>)
```




```
plt.hist(df['Population'])
```

```
(array([2457., 819., 364., 273., 91., 91., 0., 0., 91.,
       91.]),
array([ 102244., 1336256.9, 2570269.8, 3804282.7, 5038295.6,
       6272308.5, 7506321.4, 8740334.3, 9974347.2, 11208360.1,
       12442373. ]),
<BarContainer object of 10 artists>)
```




```
plt.hist(df['Metro_Population'])
```

```
(array([2639., 910., 182., 91., 182., 0., 0., 91., 91.,
       91.]),
 array([ 102244. , 1933448.4, 3764652.8, 5595857.2, 7427061.6,
        9258266. , 11089470.4, 12920674.8, 14751879.2, 16583083.6,
        18414288. ]),
<BarContainer object of 10 artists>)
```



```
plt.hist(df['Sexratio'])

(array([ 91., 0., 455., 728., 1274., 728., 182., 364., 182.,
        273.]),
 array([ 756. , 789.7, 823.4, 857.1, 890.8, 924.5, 958.2, 991.9,
        1025.6, 1059.3, 1093. ]),
<BarContainer object of 10 artists>)
```



> **b) Scatter plot**

[] ↪ 4 cells hidden

> **c) pie chart**

[] ↪ 1 cell hidden


> **d) barchart**

[] ↪ 1 cell hidden

✓ 3. Handling outliers.

```
sns.boxplot(x='Metro_Population', data=df)
```

<Axes: xlabel='Metro_Population'>




```
print(df['Metro_Population'].quantile(0.80))
```

2920067.0


```
df['Metro_Population'] = np.where(df['Metro_Population'] > 2920067, 2920067, df['Metro_Population'])
```

```
sns.boxplot(x='Metro_Population', data=df)
```

<Axes: xlabel='Metro_Population'>



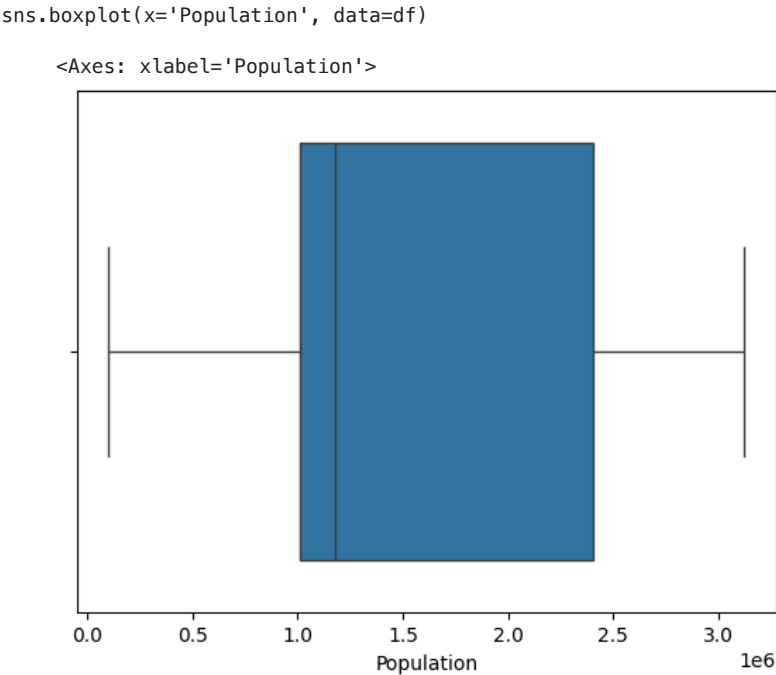
```
sns.boxplot(x='Population', data=df)
```



```
val = df['Population'].quantile(0.85)
print(val)


3124458.0

df['Population'] = np.where(df['Population'] > val, val, df['Population'])
```



```
sns.boxplot(x='Sexratio', data=df)
```

<Axes: xlabel='Sexratio'>



```
val = df['Sexratio'].quantile(0.93)
print(val)
```

```
1054.0
```


```
val_left = df['Sexratio'].quantile(0.05)
print(val_left)
```

```
850.0
```

```
df['Sexratio'] = np.where(df['Sexratio'] > val, val, df['Sexratio'])
df['Sexratio'] = np.where(df['Sexratio'] < val_left, val_left, df['Sexratio'])
```


```
sns.boxplot(x='Sexratio', data=df)
```

<Axes: xlabel='Sexratio'>



```
sns.boxplot(x='Literacy', data=df)
```

<Axes: xlabel='Literacy'>




```
val_left = df['Literacy'].quantile(0.05)  
print(val_left)
```

75.66

```
df['Literacy'] = np.where(df['Literacy'] < val_left, val_left, df['Literacy'])
```


```
sns.boxplot(x='Literacy', data=df)
```

<Axes: xlabel='Literacy'>



```
sns.boxplot(x='Unemployment', data=df)
```

<Axes: xlabel='Unemployment'>




```
val = df['Unemployment'].quantile(0.80)  
print(val)
```

12.4


```
df['Unemployment'] = np.where(df['Unemployment'] > val, val, df['Unemployment'])
```

```
sns.boxplot(x='Unemployment', data=df)
```


<Axes: xlabel='Unemployment'>



```
sns.boxplot(x='2010', data=df)
```

```
sns.boxplot(x='avg', data=df)
```



```
df.head()
```

	Rank	city	State	Population	Metro_Population	Sexratio	Literacy	Unemployment	Poverty (MPI)	crime_name	2010
0	4	Hyderabad	Andhra Pradesh	6731790	2920067	955	83.26	95.0	NaN	Acid attack	NaN
1	4	Hyderabad	Andhra Pradesh	6731790	2920067	955	83.26	95.0	NaN	Agrarian riots	NaN
2	4	Hyderabad	Andhra Pradesh	6731790	2920067	955	83.26	95.0	NaN	Arson	0.5
3	4	Hyderabad	Andhra Pradesh	6731790	2920067	955	83.26	95.0	NaN	Assault on women with intent to outrage her mo...	4.5
4	4	Hyderabad	Andhra Pradesh	6731790	2920067	955	83.26	95.0	NaN	Attempt to acid attack	NaN

Next steps:

[View recommended plots](#)

4. Handling missing values.

```
df.isna().sum()

Rank          0
city          0
State         0
Population    0
Metro_Population 0
Sexratio      0
Literacy      0
Poverty (MPI) 364
crime_name    0
2010         3285
2011         2820
2012         2820
2013         2820
2014          0
avg           0
dtype: int64
```

```
df['Poverty (MPI)'].fillna(df['Poverty (MPI)'].mean(), inplace= True)
```

```
condition = (df['2010'].isnull() & df['2011'].isnull() & df['2012'].isnull() & df['2013'].isnull())
```

```
filtered_df = df[~condition]
```

```
filtered_df
```

	Rank	city	State	Population	Metro_Population	Sexratio	Literacy	Poverty (MPI)	crime_name	2010	2011	2012	2013	2014
2	4	15	0	6731790	2920067	955	83.26	0.049	2	0.3	0.6	1.3	0.6	0.3
3	4	15	0	6731790	2920067	955	83.26	0.049	3	4.5	4.2	3.9	6.0	4.5
6	4	15	0	6731790	2920067	955	83.26	0.049	6	2.6	1.8	2.1	1.8	2.6
8	4	15	0	6731790	2920067	955	83.26	0.049	8	27.9	23.5	17.7	17.7	27.9
9	4	15	0	6731790	2920067	955	83.26	0.049	9	16.4	8.9	8.3	8.3	16.4
...
4260	78	4	16	563917	1243008	929	83.30	0.049	74	4.5	8.0	17.1	23.5	4.5
4262	78	4	16	563917	1243008	929	83.30	0.049	76	0.5	1.9	6.0	6.0	0.5
4267	78	4	16	563917	1243008	929	83.30	0.049	81	2.2	6.2	5.9	17.7	2.2
4270	78	4	16	563917	1243008	929	83.30	0.049	84	29.6	51.6	76.2	51.6	29.6
4271	78	4	16	563917	1243008	929	83.30	0.049	85	168.7	272.0	313.5	381.0	168.7

1504 rows x 15 columns

```
filtered_df.isna().sum()

Rank          0
city          0
State         0
Population    0
Metro_Population 0
Sexratio      0
Literacy      0
Poverty (MPI) 0
crime_name    0
2010         512
2011          47
2012          47
2013          47
2014          0
avg           0
dtype: int64
```

5. Handling categorical features.

6. Feature scaling.

a) Normalization

```
from sklearn.preprocessing import Normalizer
```

```
scaler = Normalizer()
scaled_data = scaler.fit_transform(filtered_df)
scaled_df = pd.DataFrame(scaled_data, columns=df.columns)
print(scaled_df.head())
```

```
-----
ValueError                                Traceback (most recent call last)
<ipython-input-65-79a0a1928bd8> in <cell line: 4>()
      2
      3 scaler = Normalizer()
----> 4 scaled_data = scaler.fit_transform(filtered_df)
      5 scaled_df = pd.DataFrame(scaled_data,
      6                          columns=df.columns)
```

```
----- 5 frames -----
/usr/local/lib/python3.10/dist-packages/sklearn/utils/validation.py in _assert_all_finite(X, allow_nan,
msg_dtype, estimator_name, input_name)
    159         "#estimators-that-handle-nan-values"
    160     )
--> 161     raise ValueError(msg_err)
    162
    163
```

ValueError: Input X contains NaN.

Normalizer does not accept missing values encoded as NaN natively. For supervised learning, you might want to consider `sklearn.ensemble.HistGradientBoostingClassifier` and `Regressor` which accept missing values encoded as